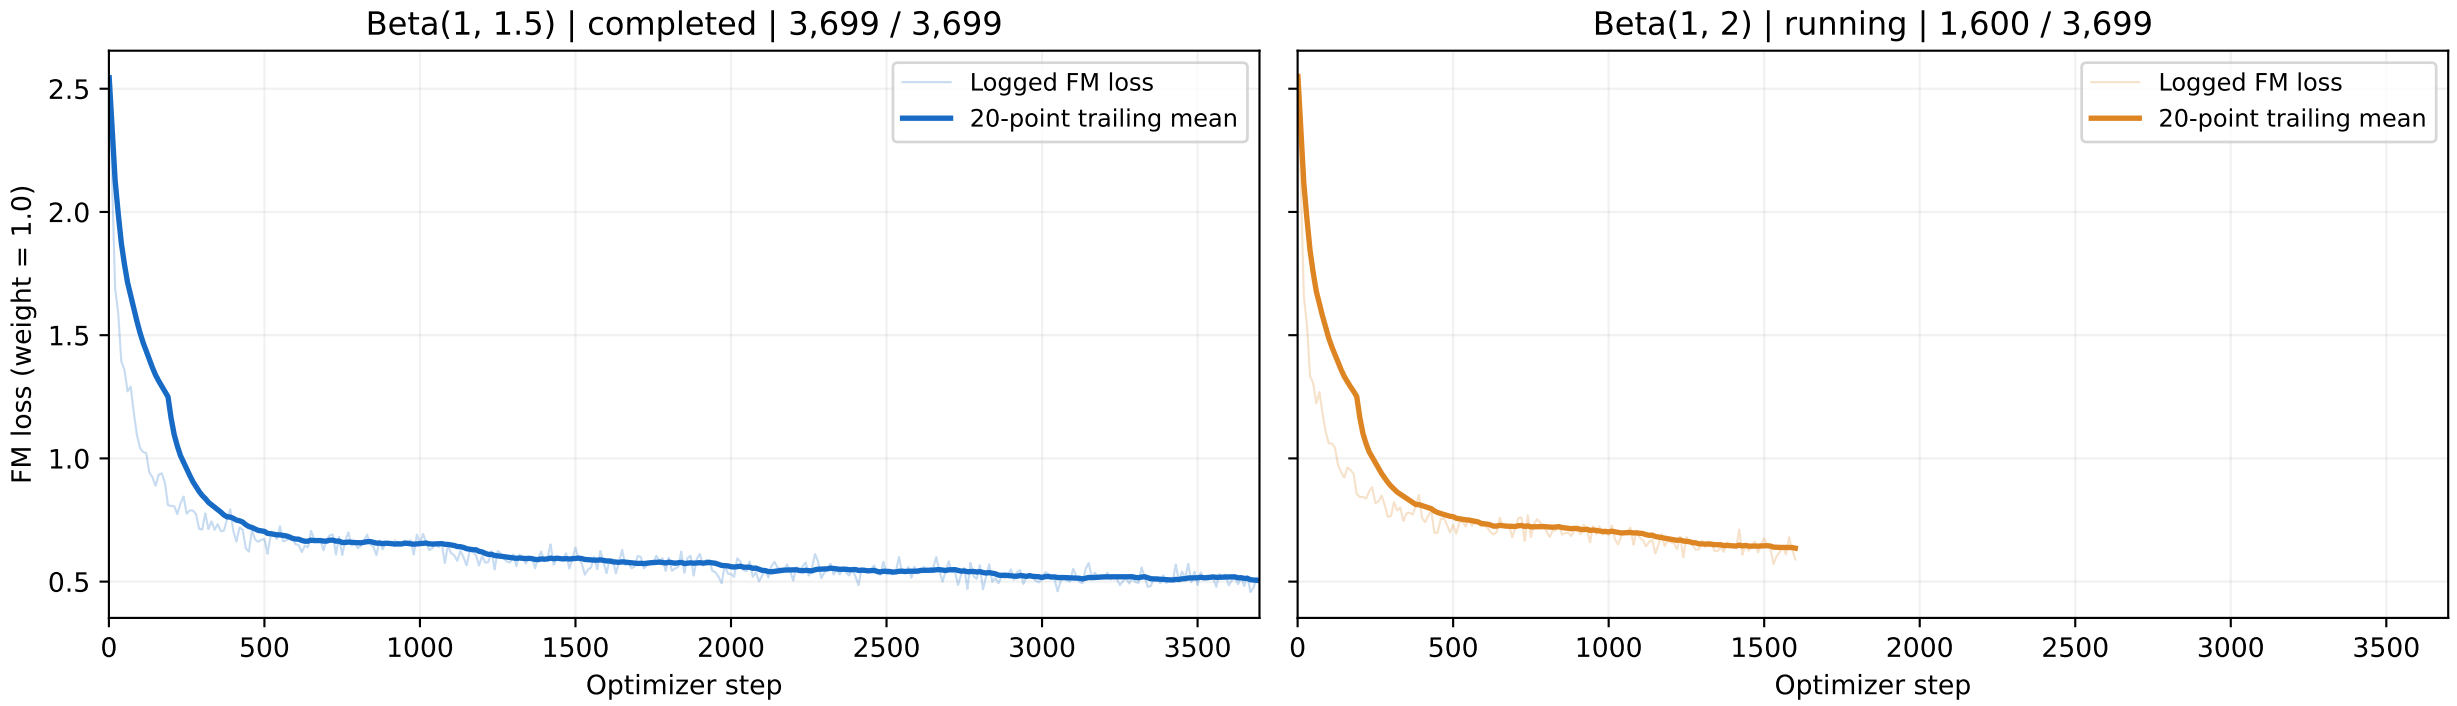
FM-only, 10 future actions x 3 dimensions | 8 GPUs | global batch 512
Beta(1, 1.5) | completed | 3,699 / 3,699
Beta(1, 2) | running | 1,600 / 3,699
Logged FM loss
20-point trailing mean
FM loss (weight = 1.0)
Optimizer step
0
500
1000
1500
2000
2500
3000
3500
0.5
1.0
1.5
2.0
2.5
Beta(1,3): queued. Different time-sampling distributions; training losses are not navigation evaluation scores.